

This project is aimed at building a predictive model to help a bank with its direct marketing. The bank has historical data for over 45,000 customers to build a predictive model to predict whether a customer will go for term deposit or not by using the attributes bank has in its database. Your goal is to study the data and build such a predictive model. You will evaluate your model by using N-fold cross-validation. Link to data for the project: <https://archive.ics.uci.edu/ml/datasets/Bank+Marketing> (<https://archive.ics.uci.edu/ml/datasets/Bank+Marketing>) You will work with the full dataset with 20 attributes. While evaluating your model, you may want to take into account cost associated with different errors.

```
In [2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

```
In [91]: bdata =pd.read_csv(r"G:/Yamuna docs/College docs/Info Ret/Project 1/bank-addit
ional-full.csv", delimiter=";",header='infer') #header='infer'
bdata.head()
```

Out[91]:

	age	job	marital	education	default	housing	loan	contact	month	day_of_week
0	56	housemaid	married	basic.4y	no	no	no	telephone	may	mon
1	57	services	married	high.school	unknown	no	no	telephone	may	mon
2	37	services	married	high.school	no	yes	no	telephone	may	mon
3	40	admin.	married	basic.6y	no	no	no	telephone	may	mon
4	56	services	married	high.school	no	no	yes	telephone	may	mon

5 rows × 21 columns



```
In [92]: #find the missing values in the dataset
print(bdata.isnull().sum())
```

```
age          0
job          0
marital      0
education    0
default      0
housing      0
loan         0
contact      0
month        0
day_of_week  0
duration     0
campaign     0
pdays       0
previous     0
poutcome     0
emp.var.rate 0
cons.price.idx 0
cons.conf.idx 0
euribor3m    0
nr.employed  0
y            0
dtype: int64
```

```
In [93]: #Dropping the duplicates
bdata = bdata.drop_duplicates()

bdata.shape
```

Out[93]: (41176, 21)

```
In [6]: #summary statistics for categorical variables
bdata.describe(include=['object'])
cate = bdata.describe(include=['object']).columns
print(cate)
bdata.describe(include=['object'])
```




```
Index(['job', 'marital', 'education', 'default', 'housing', 'loan', 'contact',
      'month', 'day_of_week', 'poutcome', 'y'],
      dtype='object')
```




Out[6]:




	job	marital	education	default	housing	loan	contact	month	day_of_week
count	41176	41176	41176	41176	41176	41176	41176	41176	41176
unique	12	4	8	3	3	3	2	10	5
top	admin.	married	university.degree	no	yes	no	cellular	may	thu
freq	10419	24921	12164	32577	21571	33938	26135	13767	8618



```
In [6]: import seaborn as sns

#Bar plots of categorical features
for feature in bdata.dtypes[bdata.describe(include=['object']).columns].index:
    sns.countplot(y=feature, data=bdata, order = bdata[feature].value_counts()
    .index)
    plt.show()
```










```
In [7]: #pd.crosstab(index=bdata["education"], columns=bdata["y"])
        bdata.education.value_counts()/bdata.education.count()
```

```
Out[7]: university.degree      0.295415
        high.school            0.231008
        basic.9y               0.146809
        professional.course    0.127259
        basic.4y               0.101418
        basic.6y               0.055639
        unknown                0.042015
        illiterate             0.000437
        Name: education, dtype: float64
```

```
In [8]: bdata.education.value_counts().plot(kind="barh")
```

```
Out[8]: <matplotlib.axes._subplots.AxesSubplot at 0x2990acb5d30>
```




```
In [9]: #pd.crosstab(index=bdata["job"], columns=bdata["y"])
bdata.job.value_counts()/bdata.job.count()
```

```
Out[9]: admin.          0.253036
blue-collar  0.224718
technician  0.163663
services    0.096343
management 0.071012
retired     0.041723
entrepreneur 0.035360
self-employed 0.034510
housemaid   0.025743
unemployed  0.024626
student     0.021250
unknown     0.008014
Name: job, dtype: float64
```

```
In [10]: bdata.job.value_counts().plot(kind="barh")
```

```
Out[10]: <matplotlib.axes._subplots.AxesSubplot at 0x2990b06ccc0>
```





```
In [11]: bdata.marital.value_counts()/bdata.marital.count()
```

```
Out[11]: married      0.605231  
single      0.280843  
divorced     0.111983  
unknown      0.001943  
Name: marital, dtype: float64
```

```
In [12]: bdata.marital.value_counts().plot(kind="barh")
```

```
Out[12]: <matplotlib.axes._subplots.AxesSubplot at 0x2990acb0470>
```




```
In [13]: bdata.default.value_counts()/bdata.default.count()
```

```
Out[13]: no          0.791165  
unknown    0.208762  
yes        0.000073  
Name: default, dtype: float64
```

```
In [14]: bdata.default.value_counts().plot(kind="barh")
```

```
Out[14]: <matplotlib.axes._subplots.AxesSubplot at 0x2990aeb5780>
```




```
In [15]: bdata.housing.value_counts()/bdata.housing.count()
```

```
Out[15]: yes          0.523873  
no          0.452084  
unknown     0.024043  
Name: housing, dtype: float64
```

```
In [16]: bdata.housing.value_counts().plot(kind="barh")
```

```
Out[16]: <matplotlib.axes._subplots.AxesSubplot at 0x2990aeb5cf8>
```




```
In [17]: bdata.loan.value_counts()/bdata.loan.count()
```

```
Out[17]: no          0.824218  
yes       0.151739  
unknown   0.024043  
Name: loan, dtype: float64
```

```
In [18]: bdata.loan.value_counts().plot(kind="barh")
```

```
Out[18]: <matplotlib.axes._subplots.AxesSubplot at 0x2990ae4be48>
```




```
In [19]: bdata.contact.value_counts()/bdata.contact.count()
```

```
Out[19]: cellular      0.634714  
telephone    0.365286  
Name: contact, dtype: float64
```

```
In [20]: bdata.contact.value_counts().plot(kind="barh")
```

```
Out[20]: <matplotlib.axes._subplots.AxesSubplot at 0x2990ae0db38>
```




```
In [21]: bdata.month.value_counts()/bdata.contact.count()
```

```
Out[21]: may      0.334345  
jul      0.174106  
aug      0.149990  
jun      0.129153  
nov      0.099573  
apr      0.063896  
oct      0.017413  
sep      0.013843  
mar      0.013260  
dec      0.004420  
Name: month, dtype: float64
```

```
In [22]: bdata.month.value_counts().plot(kind="barh")
```

```
Out[22]: <matplotlib.axes._subplots.AxesSubplot at 0x2990b0c9cf8>
```




```
In [23]: bdata.day_of_week.value_counts()/bdata.day_of_week.count()
```

```
Out[23]: thu    0.209297  
mon    0.206722  
wed    0.197542  
tue    0.196377  
fri    0.190062  
Name: day_of_week, dtype: float64
```

```
In [24]: bdata.day_of_week.value_counts().plot(kind="barh")
```

```
Out[24]: <matplotlib.axes._subplots.AxesSubplot at 0x2990b14ec88>
```




```
In [25]: bdata.poutcome.value_counts()/bdata.poutcome.count()
```

```
Out[25]: nonexistent    0.863391  
failure    0.103264  
success    0.033345  
Name: poutcome, dtype: float64
```

```
In [26]: bdata.poutcome.value_counts().plot(kind="barh")
```

```
Out[26]: <matplotlib.axes._subplots.AxesSubplot at 0x2990b1b4dd8>
```




```
In [27]: bdata.y.value_counts()/bdata.y.count()
```

```
Out[27]: no      0.887337  
yes      0.112663  
Name: y, dtype: float64
```

```
In [28]: bdata.y.value_counts().plot(kind="barh")
```

```
Out[28]: <matplotlib.axes._subplots.AxesSubplot at 0x2990b062ac8>
```



```
In [29]: #bdata.duration.value_counts()/bdata.duration.count()
```

```
In [30]: #summary statistics for quantitative variables
bdata.describe()
cont = bdata.describe().columns
print(cont)
bdata.describe()
```

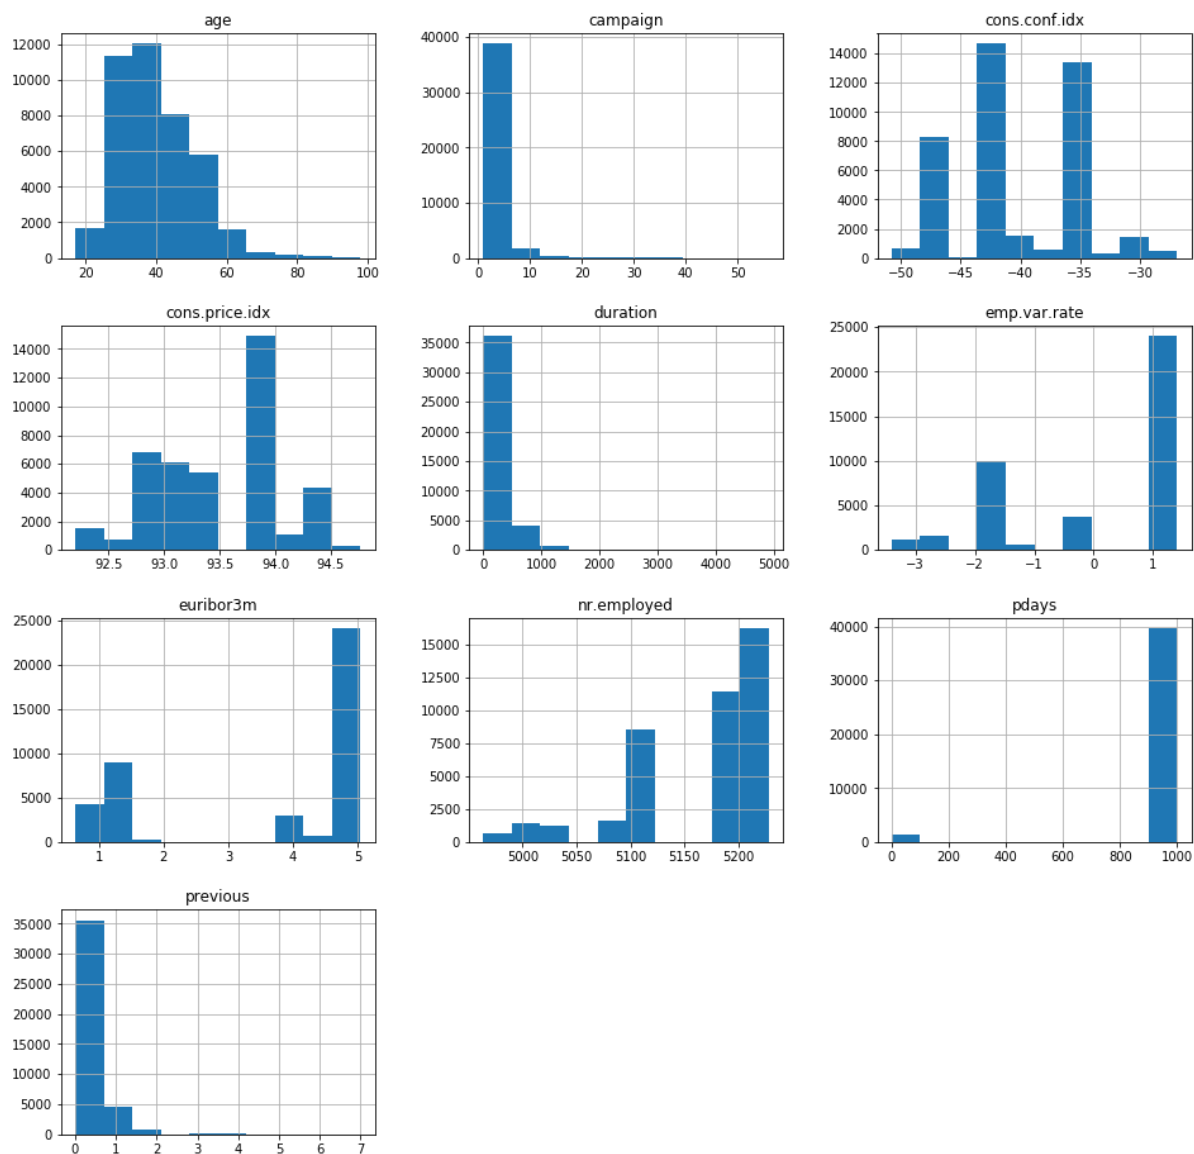
```
Index(['age', 'duration', 'campaign', 'pdays', 'previous', 'emp.var.rate',
       'cons.price.idx', 'cons.conf.idx', 'euribor3m', 'nr.employed'],
      dtype='object')
```

Out[30]:

	age	duration	campaign	pdays	previous	emp.var.rate	cons.
count	41176.00000	41176.000000	41176.000000	41176.000000	41176.000000	41176.000000	4117
mean	40.02380	258.315815	2.567879	962.464810	0.173013	0.081922	9
std	10.42068	259.305321	2.770318	186.937102	0.494964	1.570883	
min	17.00000	0.000000	1.000000	0.000000	0.000000	-3.400000	9
25%	32.00000	102.000000	1.000000	999.000000	0.000000	-1.800000	9
50%	38.00000	180.000000	2.000000	999.000000	0.000000	1.100000	9
75%	47.00000	319.000000	3.000000	999.000000	0.000000	1.400000	9
max	98.00000	4918.000000	56.000000	999.000000	7.000000	1.400000	9

In [195]: *#Histogram Visualizing Distribution of quantitative variables*

```
cont_histogram = bdata.hist(column=cont, figsize = (16,16))
```



```
In [7]: #Statistical Summary
bdata.describe()
```

Out[7]:

	age	duration	campaign	pdays	previous	emp.var.rate	cons.
count	41176.00000	41176.000000	41176.000000	41176.000000	41176.000000	41176.000000	4117
mean	40.02380	258.315815	2.567879	962.464810	0.173013	0.081922	9
std	10.42068	259.305321	2.770318	186.937102	0.494964	1.570883	
min	17.00000	0.000000	1.000000	0.000000	0.000000	-3.400000	9
25%	32.00000	102.000000	1.000000	999.000000	0.000000	-1.800000	9
50%	38.00000	180.000000	2.000000	999.000000	0.000000	1.100000	9
75%	47.00000	319.000000	3.000000	999.000000	0.000000	1.400000	9
max	98.00000	4918.000000	56.000000	999.000000	7.000000	1.400000	9


```
In [32]: #find the unknown values in the dataset
unknown_values = ["unknown", "na", "--"]
bdata = pd.read_csv(r"G:/Yamuna docs/College docs/Info Ret/Project 1/bank-addi
tional-full.csv", delimiter=";", header='infer', na_values = unknown_values)
print(bdata.isnull().sum())
bdata.isnull().sum().sum()
```

```
age          0
job          330
marital      80
education    1731
default      8597
housing      990
loan         990
contact      0
month        0
day_of_week  0
duration     0
campaign     0
pdays       0
previous     0
poutcome    0
emp.var.rate 0
cons.price.idx 0
cons.conf.idx 0
euribor3m    0
nr.employed  0
y            0
dtype: int64
```

Out[32]: 12718


```
In [292]: import missingno as msno  
msno.matrix(bdata)
```

```
Out[292]: <matplotlib.axes._subplots.AxesSubplot at 0x133d1220828>
```



```
In [215]: bdata = bdata[bdata.job != 'unknown']
bdata = bdata[bdata.marital != 'unknown']

bdata = bdata[bdata.education != 'unknown']
bdata = bdata[bdata.education != 'illiterate']

#Dropping the unknown housing loan status
bdata = bdata[bdata.housing != 'unknown']

#Dropping the unknown personal loan status
bdata = bdata[bdata.loan != 'unknown']

#Deleting the 'default' column
del bdata['default']

#Deleting the 'duration' column
del bdata['duration']
bdata.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 41170 entries, 0 to 41187
Data columns (total 20 columns):
age                41170 non-null int64
job                40840 non-null object
marital            41090 non-null object
education          39439 non-null object
housing            40180 non-null object
loan               40180 non-null object
contact            41170 non-null object
month              41170 non-null object
day_of_week        41170 non-null object
duration           41170 non-null int64
campaign           41170 non-null int64
pdays            41170 non-null int64
previous           41170 non-null int64
poutcome          41170 non-null object
emp.var.rate       41170 non-null float64
cons.price.idx     41170 non-null float64
cons.conf.idx      41170 non-null float64
euribor3m          41170 non-null float64
nr.employed        41170 non-null float64
y                  41170 non-null object
dtypes: float64(5), int64(5), object(10)
memory usage: 6.6+ MB
```

```
In [94]: #combining spares categories
bdata.job.replace(['entrepreneur', 'self-employed'], 'self-employed', inplace=True)

bdata.job.replace(['admin.', 'management'], 'administration_management', inplace=True)

bdata.job.replace(['blue-collar', 'technician'], 'blue-collar', inplace=True)

bdata.job.replace(['retired', 'unemployed'], 'no_active_income', inplace=True)

bdata.job.replace(['services', 'housemaid'], 'services', inplace=True)

bdata.marital.replace(['single', 'divorced'], 'single', inplace=True)

bdata.education.replace(['basic.9y', 'basic.6y', 'basic.4y'], 'basic_school',
inplace=True)
```

```
In [58]: catanz = bdata.loc[( bdata['pdays'] == 999) & (bdata['poutcome'] != 'nonexistent') ]

#Counting the values
catanz.poutcome.value_counts()
```

```
Out[58]: failure    4110
Name: poutcome, dtype: int64
```

```
In [59]: ind_999 = bdata.loc[(bdata['pdays'] == 999) & (bdata['poutcome'] != 'nonexistent')]['pdays'].index.values
```

```
#Assigning NaNs instead of '999'
```

```
bdata.loc[ind_999, 'pdays'] = np.nan
```

```
#Checking if the NaNs were assigned correctly
```

```
bdata.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 41176 entries, 0 to 41187
Data columns (total 21 columns):
age                41176 non-null int64
job                41176 non-null object
marital            41176 non-null object
education          41176 non-null object
default            41176 non-null object
housing            41176 non-null object
loan               41176 non-null object
contact            41176 non-null object
month              41176 non-null object
day_of_week        41176 non-null object
duration           41176 non-null int64
campaign           41176 non-null int64
pdays             37066 non-null float64
previous           41176 non-null int64
poutcome           41176 non-null object
emp.var.rate       41176 non-null float64
cons.price.idx     41176 non-null float64
cons.conf.idx      41176 non-null float64
euribor3m          41176 non-null float64
nr.employed        41176 non-null float64
y                  41176 non-null object
dtypes: float64(6), int64(4), object(11)
memory usage: 8.2+ MB
```

In [187]: *#Dropping NAs from the dataset*

```
bdata = bdata.dropna()
bdata.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 0 entries
Data columns (total 20 columns):
age                0 non-null float64
job                0 non-null float64
marital            0 non-null float64
education          0 non-null float64
housing            0 non-null float64
loan               0 non-null float64
contact            0 non-null float64
month              0 non-null float64
day_of_week        0 non-null float64
duration           0 non-null int64
campaign           0 non-null float64
pdays            0 non-null float64
previous           0 non-null float64
poutcome          0 non-null float64
emp.var.rate       0 non-null float64
cons.price.idx     0 non-null float64
cons.conf.idx      0 non-null float64
euribor3m          0 non-null float64
nr.employed        0 non-null float64
y                  0 non-null int32
dtypes: float64(18), int32(1), int64(1)
memory usage: 0.0 bytes
```

In [34]: *#correlation matrix*


```
bdata.corr()
```

Out[34]:

	age	duration	campaign	pdays	previous	emp.var.rate	cons.price.idx
age	1.000000	-0.000866	0.004594	-0.034369	0.024365	-0.000371	0.000857
duration	-0.000866	1.000000	-0.071699	-0.047577	0.020640	-0.027968	0.005312
campaign	0.004594	-0.071699	1.000000	0.052584	-0.079141	0.150754	0.127836
pdays	-0.034369	-0.047577	0.052584	1.000000	-0.587514	0.271004	0.078889
previous	0.024365	0.020640	-0.079141	-0.587514	1.000000	-0.420489	-0.203130
emp.var.rate	-0.000371	-0.027968	0.150754	0.271004	-0.420489	1.000000	0.775334
cons.price.idx	0.000857	0.005312	0.127836	0.078889	-0.203130	0.775334	1.000000
cons.conf.idx	0.129372	-0.008173	-0.013733	-0.091342	-0.050936	0.196041	0.058986
euribor3m	0.010767	-0.032897	0.135133	0.296899	-0.454494	0.972245	0.688230
nr.employed	-0.017725	-0.044703	0.144095	0.372605	-0.501333	0.906970	0.522034

```
In [35]: import seaborn as sns
#visualise correlation
sns.heatmap(bdata.corr(), )
```

Out[35]: <matplotlib.axes._subplots.AxesSubplot at 0x2990b063080>



In []: From the above analysis, data **is** non-linear **and** its an asymmetric **type**. The feature selection will **not** depend on the correlation matrix.

```
In [36]: import pandas as pd
#Class Distribution
bdata['y'].value_counts()
#bdata.groupby('y').count()#.toPandas()
```

Out[36]: no 36548
yes 4640
Name: y, dtype: int64

```
In [767]: #features = bdata[bdata.columns.difference(['y'])]
#Labels = bdata['y']
#Labels
#bdata.fillna(0, inplace=True)
#bad_chars = [';', ':', '!', "*", '.']
#bdata = ''.join(i for i in bdata if not i in bad_chars)
#bdata=bdata.replace('\.', '', regex=True).astype(object)
```

```
In [95]: #bdata.drop(['duration', 'contact', 'month', 'default', 'day_of_week', 'pdays', ], ax
is=1, inplace=True)
bdata.drop(['default'], axis=1, inplace=True)
```

```
In [96]: #converting categorical features into numeric values
from sklearn import preprocessing
from sklearn.preprocessing import LabelEncoder
le = preprocessing.LabelEncoder()
```

```
In [77]: bdata.job = bdata.job.replace(['<'], '0', inplace=True)
bdata.marital = bdata.marital.replace(['<'], '0', inplace=True)
bdata.education = bdata.education.replace(['<'], '0', inplace=True)
bdata.housing = bdata.housing.replace(['<'], '0', inplace=True)
bdata.loan = bdata.loan.replace(['<'], '0', inplace=True)
bdata.poutcome = bdata.poutcome.replace(['<'], '0', inplace=True)
#bdata.default = bdata.default.replace(['<'], '0', inplace=True)
```

```
In [97]: bdata.job = le.fit_transform(bdata.job)
bdata.marital = le.fit_transform(bdata.marital)
bdata.education = le.fit_transform(bdata.education)
bdata.housing = le.fit_transform(bdata.housing)
bdata.loan = le.fit_transform(bdata.loan)
bdata.poutcome = le.fit_transform(bdata.poutcome)
bdata.y = le.fit_transform(bdata.y)
bdata.day_of_week = le.fit_transform(bdata.day_of_week)
#bdata.default = le.fit_transform(bdata.default)
bdata.contact = le.fit_transform(bdata.contact)
bdata.month = le.fit_transform(bdata.month)

bdata.head()
```

Out[97]:

	age	job	marital	education	housing	loan	contact	month	day_of_week	duration	campaign
0	56	4	0	0	0	0	1	6	1	261	
1	57	4	0	1	0	0	1	6	1	149	
2	37	4	0	1	2	0	1	6	1	226	
3	40	0	0	0	0	0	1	6	1	151	
4	56	4	0	1	0	2	1	6	1	307	



```
In [98]: bdata.shape
```

Out[98]: (41176, 20)

```
In [99]: num_features = ['age', 'job', 'marital', 'education', 'housing', 'loan',
                        'contact', 'month', 'day_of_week', 'campaign',
                        'previous', 'poutcome', 'emp.var.rate', 'cons.price.idx',
                        'cons.conf.idx', 'euribor3m', 'nr.employed']

scaled_features = {}
for each in num_features:
    mean, std = bdata[each].mean(), bdata[each].std()
    scaled_features[each] = [mean, std]
    bdata.loc[:, each] = (bdata[each] - mean)/std
print(pd.DataFrame(scaled_features))
```

	age	job	marital	education	housing	loan	contact \
0	40.02380	1.373373	0.396712	2.005391	1.071789	0.327521	0.365286
1	10.42068	1.472977	0.493177	1.770190	0.985305	0.723700	0.481516

	month	day_of_week	campaign	previous	poutcome	emp.var.rate \
0	4.231033	2.004614	2.567879	0.173013	0.930081	0.081922
1	2.319973	1.397692	2.770318	0.494964	0.362937	1.570883

	cons.price.idx	cons.conf.idx	euribor3m	nr.employed
0	93.575720	-40.502863	3.621293	5167.034870
1	0.578839	4.627860	1.734437	72.251364

```
In [100]: X = bdata.iloc[:,0:19]
X.columns
```

```
Out[100]: Index(['age', 'job', 'marital', 'education', 'housing', 'loan', 'contact',
                'month', 'day_of_week', 'duration', 'campaign', 'pdays', 'previous',
                'poutcome', 'emp.var.rate', 'cons.price.idx', 'cons.conf.idx',
                'euribor3m', 'nr.employed'],
                dtype='object')
```

```
In [101]: y = bdata.iloc[:,19]
```

```
In [102]: y.value_counts()
```

```
Out[102]: 0    36537
          1     4639
          Name: y, dtype: int64
```

```
In [103]: from sklearn import model_selection
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegressionCV, LogisticRegression
from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
from sklearn.metrics import confusion_matrix, accuracy_score
k_fold = KFold(n_splits=10, shuffle=True, random_state=0)
x_train, x_test, y_train, y_test = model_selection.train_test_split(X, y, test
_size=0.2, random_state=0)
```



```
In [104]: from sklearn.preprocessing import StandardScaler
          sc= StandardScaler()
          x_train = sc.fit_transform(x_train)
          x_test = sc.transform(x_test)
```

```
In [105]: x_train.shape, y_train.shape
```

```
Out[105]: ((32940, 19), (32940,))
```

```
In [106]: x_test.shape, y_test.shape
```

```
Out[106]: ((8236, 19), (8236,))
```

```
In [107]: model=LogisticRegression(penalty='l2', max_iter=10)
          modellr=model.fit(x_train, y_train)
          modellr
```

C:\ProgramData\Anaconda3\envs\py36\lib\site-packages\sklearn\linear_model\logistic.py:432: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.
FutureWarning)

```
Out[107]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                             intercept_scaling=1, l1_ratio=None, max_iter=10,
                             multi_class='warn', n_jobs=None, penalty='l2',
                             random_state=None, solver='warn', tol=0.0001, verbose=0,
                             warm_start=False)
```

```
In [108]: prediction=model.predict(x_test) #predict
```

```
In [109]: from sklearn.metrics import accuracy_score
          print('accuracy_score',accuracy_score(y_test, prediction))
```

```
accuracy_score 0.9091792132102963
```

```
In [110]: accuracy = cross_val_score(modellr, x_train, y_train, scoring='accuracy', cv =
k_fold, n_jobs=1)
print('Accuracy of each fold:',accuracy)
#get the mean of each fold
Logic = accuracy.mean() * 100
print("Accuracy of Model with Cross Validation:",Logic)
```

```
C:\ProgramData\Anaconda3\envs\py36\lib\site-packages\sklearn\linear_model\log
istic.py:432: FutureWarning: Default solver will be changed to 'lbfgs' in 0.2
2. Specify a solver to silence this warning.
```

```
FutureWarning)
```

```
C:\ProgramData\Anaconda3\envs\py36\lib\site-packages\sklearn\linear_model\log
istic.py:432: FutureWarning: Default solver will be changed to 'lbfgs' in 0.2
2. Specify a solver to silence this warning.
```

```
FutureWarning)
```

```
C:\ProgramData\Anaconda3\envs\py36\lib\site-packages\sklearn\linear_model\log
istic.py:432: FutureWarning: Default solver will be changed to 'lbfgs' in 0.2
2. Specify a solver to silence this warning.
```

```
FutureWarning)
```

```
C:\ProgramData\Anaconda3\envs\py36\lib\site-packages\sklearn\linear_model\log
istic.py:432: FutureWarning: Default solver will be changed to 'lbfgs' in 0.2
2. Specify a solver to silence this warning.
```

```
FutureWarning)
```

```
C:\ProgramData\Anaconda3\envs\py36\lib\site-packages\sklearn\linear_model\log
istic.py:432: FutureWarning: Default solver will be changed to 'lbfgs' in 0.2
2. Specify a solver to silence this warning.
```

```
FutureWarning)
```

```
C:\ProgramData\Anaconda3\envs\py36\lib\site-packages\sklearn\linear_model\log
istic.py:432: FutureWarning: Default solver will be changed to 'lbfgs' in 0.2
2. Specify a solver to silence this warning.
```

```
FutureWarning)
```

```
C:\ProgramData\Anaconda3\envs\py36\lib\site-packages\sklearn\linear_model\log
istic.py:432: FutureWarning: Default solver will be changed to 'lbfgs' in 0.2
2. Specify a solver to silence this warning.
```

```
FutureWarning)
```

```
C:\ProgramData\Anaconda3\envs\py36\lib\site-packages\sklearn\linear_model\log
istic.py:432: FutureWarning: Default solver will be changed to 'lbfgs' in 0.2
2. Specify a solver to silence this warning.
```

```
FutureWarning)
```

```
C:\ProgramData\Anaconda3\envs\py36\lib\site-packages\sklearn\linear_model\log
istic.py:432: FutureWarning: Default solver will be changed to 'lbfgs' in 0.2
2. Specify a solver to silence this warning.
```

```
FutureWarning)
```

```
C:\ProgramData\Anaconda3\envs\py36\lib\site-packages\sklearn\linear_model\log
istic.py:432: FutureWarning: Default solver will be changed to 'lbfgs' in 0.2
2. Specify a solver to silence this warning.
```

```
FutureWarning)
```

```
Accuracy of each fold: [0.91044323 0.9058895 0.91044323 0.90892532 0.9119611
4 0.91469338
```

```
0.9058895 0.91165756 0.91621129 0.91105039]
```

```
Accuracy of Model with Cross Validation: 91.07164541590771
```

```
In [111]: from sklearn.metrics import confusion_matrix, classification_report
confusion_matrix = confusion_matrix(y_test, prediction)
print(confusion_matrix)
```

```
[[7102  167]
 [ 581  386]]
```

```
In [112]: print(classification_report(y_test, prediction))
```

	precision	recall	f1-score	support
0	0.92	0.98	0.95	7269
1	0.70	0.40	0.51	967
accuracy			0.91	8236
macro avg	0.81	0.69	0.73	8236
weighted avg	0.90	0.91	0.90	8236


```
In [113]: from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score

from sklearn import metrics
y_pred_proba1 = model.predict_proba(x_test)[::,1]
fpr, tpr, _ = metrics.roc_curve(y_test, y_pred_proba1)
auc = metrics.roc_auc_score(y_test, y_pred_proba1)
result1 = roc_auc_score(y_test, prediction)
print('AUC',auc)
```

```
AUC 0.9334926988758058
```

```
In [114]: #Obtaining the ROC score
roc_auc = roc_auc_score(y_test, y_pred_proba)
print(roc_auc)
#Obtaining false and true positives & thresholds
fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba)
plt.plot(fpr, tpr, label='LogisticRegression (AUC = %0.03f)' % roc_auc)
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC curve for Logistic regression')
plt.legend(loc="lower right")
plt.show()
```

0.9334926988758058



```
In [115]: from sklearn.tree import DecisionTreeClassifier
tree = DecisionTreeClassifier(criterion="gini", max_depth=7)
model = tree.fit(x_train, y_train)
model
```

```
Out[115]: DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=7,
max_features=None, max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, presort=False,
random_state=None, splitter='best')
```

```
In [116]: dec_pred=tree.predict(x_test)
print('Accuracy_score', accuracy_score(y_test, dec_pred))
```

Accuracy_score 0.9082078678970374

```
In [117]: #mat = confusion_matrix([y_test],[dec_pred],labels=['no','yes'])
# print(mat)
#y_test = label_binarize(y_test,classes=['no','yes'])
#y_pred = label_binarize(y_pred,classes=['no','yes'])
#print("Precision: ",round(precision_score(y_test,dec_pred),2),"Recall: ",round(recall_score(y_test,dec_pred),2))
#confusion_matrix = pd.crosstab(['y_test'], bdata['dec_pred'], rownames=['Actual'], colnames=['Predicted'])
#print (confusion_matrix)
```

```
In [118]: accuracy = cross_val_score(model, x_train, y_train, scoring='accuracy', cv = k
_fold, n_jobs=1)
print('Accuracy of each fold:',accuracy)
#get the mean of each fold
Decision = accuracy.mean()*100
print("Accuracy of Model with Cross Validation:",accuracy.mean() * 100)
```

Accuracy of each fold: [0.91256831 0.91256831 0.90801457 0.91256831 0.91013965 0.91378264 0.91621129 0.91256831 0.91833637 0.91469338]
Accuracy of Model with Cross Validation: 91.31451123254402

```
In [119]: print(classification_report(y_test,dec_pred))
```


	precision	recall	f1-score	support
0	0.93	0.97	0.95	7269
1	0.65	0.48	0.55	967
accuracy			0.91	8236
macro avg	0.79	0.72	0.75	8236
weighted avg	0.90	0.91	0.90	8236

```
In [120]: from sklearn import metrics
dec_pred_proba = model.predict_proba(x_test)[:,:1]
fpr, tpr, _ = metrics.roc_curve(y_test, dec_pred_proba)
auc = metrics.roc_auc_score(y_test, dec_pred_proba)
resultdec = roc_auc_score(y_test, dec_pred_proba)
print(auc)
```

0.9270856406979932

```
In [121]: #Obtaining the ROC score
roc_auc = roc_auc_score(y_test, dec_pred_proba)
print(roc_auc)
#Obtaining false and true positives & thresholds
fpr, tpr, thresholds = roc_curve(y_test, dec_pred_proba)
plt.plot(fpr, tpr, label='DecisionTreeClassifier (AUC = %0.03f)' % roc_auc)
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC curve for DecisionTreeClassifier')
plt.legend(loc='lower right')
plt.show()
```

0.9270856406979932



```
In [122]: from sklearn.ensemble import RandomForestClassifier

#Create a Gaussian Classifier
clf=RandomForestClassifier(n_estimators=100)

#Train the model using the training sets y_pred=clf.predict(X_test)
clf.fit(x_train,y_train)

rd_pred=clf.predict(x_test)
```

```
In [123]: print("Accuracy:",metrics.accuracy_score(y_test, rd_pred))
```

Accuracy: 0.9116075764934434

```
In [124]: accuracy = cross_val_score(clf, x_train, y_train, scoring='accuracy', cv = k_fold, n_jobs=1)
print('Accuracy of each fold:',accuracy)
#get the mean of each fold
Random = accuracy.mean()*100
print("Accuracy of Model with Cross Validation:",Random)
```

Accuracy of each fold: [0.91560413 0.91621129 0.91165756 0.91044323 0.91621129 0.91347905 0.91499696 0.9177292 0.92380085 0.91530055]
 Accuracy of Model with Cross Validation: 91.55434122647237

```
In [125]: print(classification_report(y_test,rd_pred))
```

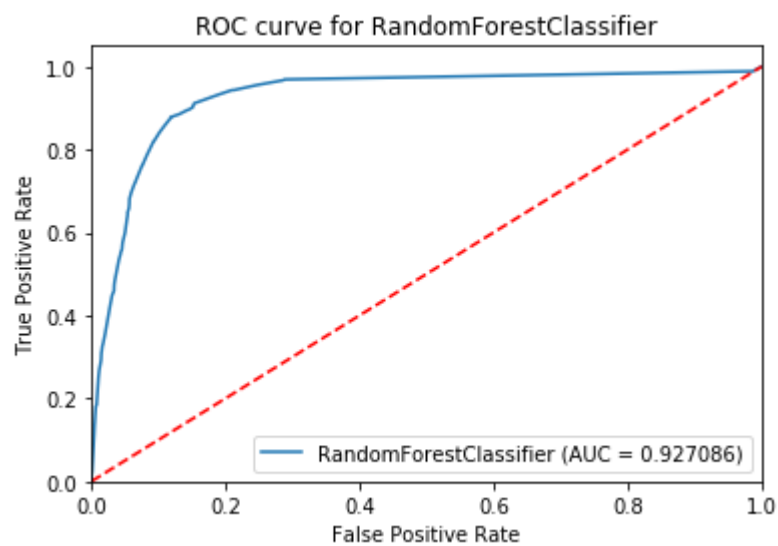
	precision	recall	f1-score	support
0	0.93	0.97	0.95	7269
1	0.67	0.49	0.57	967
accuracy			0.91	8236
macro avg	0.80	0.73	0.76	8236
weighted avg	0.90	0.91	0.91	8236

```
In [126]: from sklearn import metrics
rd_pred_proba = model.predict_proba(x_test)[::,1]
fpr, tpr, _ = metrics.roc_curve(y_test, rd_pred_proba)
auc = metrics.roc_auc_score(y_test, rd_pred_proba)
resultrd = roc_auc_score(y_test, rd_pred)
print(auc)
```

0.9270856406979932

```
In [127]: #Obtaining the ROC score
roc_auc = roc_auc_score(y_test, rd_pred_proba)
print(roc_auc)
#Obtaining false and true positives & thresholds
fpr, tpr, thresholds = roc_curve(y_test, rd_pred_proba)
plt.plot(fpr, tpr, label='RandomForestClassifier (AUC = %f)' % roc_auc)
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC curve for RandomForestClassifier')
plt.legend(loc='lower right')
plt.show()
```

0.9270856406979932




```
In [49]: from sklearn import model_selection
from sklearn.neighbors import KNeighborsClassifier

#Neighbors
neighbors = np.arange(0,25)

#Create empty list that will hold cv scores
cv_scores = []

#Perform 10-fold cross validation on training set for odd values of k:
for k in neighbors:
    k_value = k+1
    knn = KNeighborsClassifier(n_neighbors = k_value, weights='uniform', p=2,
metric='euclidean')
    kfold = model_selection.KFold(n_splits=10, random_state=123)
    scores = model_selection.cross_val_score(knn, x_train, y_train, cv=kfold,
scoring='accuracy')
    cv_scores.append(scores.mean()*100)
    print("k=%d %0.2f (+/- %0.2f)" % (k_value, scores.mean()*100, scores.std()
*100))

optimal_k = neighbors[cv_scores.index(max(cv_scores))]
print ("The optimal number of neighbors is %d with %0.1f%%" % (optimal_k, cv_s
cores[optimal_k]))


plt.plot(neighbors, cv_scores)
plt.xlabel('Number of Neighbors K')
plt.ylabel('Train Accuracy')
plt.show()
```

```

k=1 84.70 (+/- 0.45)
k=2 89.21 (+/- 0.42)
k=3 88.32 (+/- 0.41)
k=4 89.46 (+/- 0.40)
k=5 89.28 (+/- 0.37)
k=6 89.62 (+/- 0.27)
k=7 89.48 (+/- 0.38)
k=8 89.66 (+/- 0.34)
k=9 89.65 (+/- 0.43)
k=10 89.80 (+/- 0.39)
k=11 89.78 (+/- 0.46)
k=12 89.87 (+/- 0.43)
k=13 89.84 (+/- 0.38)
k=14 89.88 (+/- 0.34)
k=15 89.87 (+/- 0.33)
k=16 89.92 (+/- 0.32)
k=17 89.89 (+/- 0.39)
k=18 89.89 (+/- 0.35)
k=19 89.88 (+/- 0.38)
k=20 89.94 (+/- 0.31)
k=21 89.91 (+/- 0.40)
k=22 89.91 (+/- 0.34)
k=23 89.94 (+/- 0.35)
k=24 90.00 (+/- 0.35)
k=25 90.02 (+/- 0.39)

```

The optimal number of neighbors is 24 with 90.0%



```

In [128]: from sklearn.neighbors import KNeighborsClassifier
knn = KNeighborsClassifier(n_neighbors=22)
knn.fit(x_train, y_train)
knnpred = knn.predict(x_test)

#print(confusion_matrix(y_test, knnpred))
print("Accuracy:", metrics.accuracy_score(y_test, knnpred))

```

Accuracy: 0.9008013598834386

```
In [129]: accuracy = cross_val_score(knn, x_train, y_train, scoring='accuracy', cv = k_f
old, n_jobs=1)
print('Accuracy of each fold:',accuracy)
#get the mean of each fold
KNN = accuracy.mean()*100
print("Accuracy of Model with Cross Validation:",KNN)
```

```
Accuracy of each fold: [0.90315726 0.89678203 0.90680024 0.89860352 0.9013357
6 0.91105039
0.90497875 0.90680024 0.9143898 0.9058895 ]
Accuracy of Model with Cross Validation: 90.49787492410442
```

```
In [130]: #confusion_matrix(y_test, knnpred)
print(classification_report(y_test, knnpred))
```


	precision	recall	f1-score	support
0	0.91	0.98	0.95	7269
1	0.69	0.29	0.40	967
accuracy			0.90	8236
macro avg	0.80	0.63	0.67	8236
weighted avg	0.89	0.90	0.88	8236

```
In [131]: knn_pred_proba = model.predict_proba(x_test)[::,1]
fpr, tpr, _ = metrics.roc_curve(y_test, knn_pred_proba)
auc = metrics.roc_auc_score(y_test, knn_pred_proba)
resultknn = roc_auc_score(y_test, knnpred)
print(auc)
```

```
0.9270856406979932
```

```
In [132]: #Obtaining the ROC score
roc_auc = roc_auc_score(y_test, knn_pred_proba)
print(roc_auc)
#Obtaining false and true positives & thresholds
fpr, tpr, thresholds = roc_curve(y_test, knn_pred_proba)
plt.plot(fpr, tpr, label='KNN (AUC = %f)' % roc_auc)
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC curve for KNN Classifier')
plt.legend(loc='lower right')
plt.show()
```

0.9270856406979932



```
In [133]: from sklearn.naive_bayes import GaussianNB
gaussiannb= GaussianNB()
gaussiannb.fit(x_train, y_train)
gaussiannbpred = gaussiannb.predict(x_test)
#gauss = gaussiannb.predict(x_test)

print(accuracy_score(y_test, gaussiannbpred ))
#print('accuracy_score',round(accuracy_score(y_test, gaussiannbpred),2)*100)
#GAUSIAN = (cross_val_score(gaussiannb, x_train, y_train, cv=k_fold, n_jobs=1,
scoring = 'accuracy').mean())
```

0.8453132588635259

```
In [134]: accuracy = cross_val_score(gaussiannb, x_train, y_train, scoring='accuracy', c
v = k_fold, n_jobs=1)
print('Accuracy of each fold:',accuracy)
#get the mean of each fold
GAUSIAN = accuracy.mean()*100
print("Accuracy of Model with Cross Validation:",accuracy.mean() * 100)
```

Accuracy of each fold: [0.83333333 0.84031573 0.83849423 0.83697632 0.8503339
4 0.84669095
0.83758349 0.84092289 0.83515483 0.83151184]
Accuracy of Model with Cross Validation: 83.91317547055252

```
In [135]: print(classification_report(y_test, gaussiannbpred))
```


	precision	recall	f1-score	support
0	0.94	0.88	0.91	7269
1	0.39	0.59	0.47	967
accuracy			0.85	8236
macro avg	0.67	0.74	0.69	8236
weighted avg	0.88	0.85	0.86	8236

```
In [136]: gauss_pred_proba = model.predict_proba(x_test)[:,:1]
fpr1, tpr1, _ = metrics.roc_curve(y_test, gauss_pred_proba)
auc = metrics.roc_auc_score(y_test, gauss_pred_proba)
resultgauss = roc_auc_score(y_test, gaussiannbpred)
print(auc)
```

0.9270856406979932

```
In [137]: #Obtaining the ROC score
roc_auc = roc_auc_score(y_test, gauss_pred_proba)
print('AUC',roc_auc)
#Obtaining false and true positives & thresholds
fpr1, tpr1, thresholds = roc_curve(y_test, gauss_pred_proba)
plt.plot(fpr1, tpr1, label='GaussianNB (AUC = %0.03f)' % roc_auc)
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC curve for GaussianNB Classifier')
plt.legend(loc='lower right')
plt.show()
```

AUC 0.9270856406979932



```
In [138]: from sklearn.ensemble import GradientBoostingClassifier
gbk = GradientBoostingClassifier()
gbk.fit(x_train, y_train)
gbkpred = gbk.predict(x_test)
#print(confusion_matrix(y_test, gbkpred ))
print('accuracy_score', (accuracy_score(y_test, gbkpred)))
#GradientBoosting = (cross_val_score(gbk, x_train, y_train, cv=k_fold, n_jobs=
1, scoring = 'accuracy').mean()))
```

accuracy_score 0.91452161243322

```
In [139]: accuracy = cross_val_score(gbk, x_train, y_train, scoring='accuracy', cv = k_fold, n_jobs=1)
print('Accuracy of each fold:',accuracy)
GradientBoosting = accuracy.mean() * 100
#get the mean of each fold
print("Accuracy of Model with Cross Validation:",GradientBoosting)
```

Accuracy of each fold: [0.92076503 0.91590771 0.91317547 0.91742562 0.91985428 0.91863995 0.91317547 0.91712204 0.92380085 0.91530055]
Accuracy of Model with Cross Validation: 91.75166970248938

```
In [140]: print(classification_report(y_test, gbkpred))
```


	precision	recall	f1-score	support
0	0.94	0.97	0.95	7269
1	0.68	0.52	0.59	967
accuracy			0.91	8236
macro avg	0.81	0.75	0.77	8236
weighted avg	0.91	0.91	0.91	8236

```
In [141]: gbk_pred_proba = model.predict_proba(x_test)[:,:1]
fpr, tpr, _ = metrics.roc_curve(y_test, gbk_pred_proba)
auc = metrics.roc_auc_score(y_test, gbk_pred_proba)
resultgbk = roc_auc_score(y_test, gbkpred)
print(auc)
```

0.9270856406979932

```
In [142]: #Obtaining the ROC score
roc_auc = roc_auc_score(y_test, gbk_pred_proba)
print('AUC',roc_auc)
#Obtaining false and true positives & thresholds
fpr, tpr, thresholds = roc_curve(y_test, gbk_pred_proba)
plt.plot(fpr, tpr, label='GradientBooting (AUC = %0.03f)' % roc_auc)
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC curve for GradientBoostingClassifier')
plt.legend(loc='lower right')
plt.show()
```

AUC 0.9270856406979932



```
In [144]: models = pd.DataFrame({
            'Models': ['Logistic Model', 'Decision Tree Classifier', 'Random
Forest Classifier', 'K-Near Neighbors', 'Gaussian NB', 'Gradient Boosting'],
            'Accuracy': [Logic, Decision, Random, KNN, GAUSIAN, GradientB
oosting],
            'AUC' : [resultl, resultdec, resultrd, resultknn, resultgauss,
resultgbk] })

models.sort_values(by='Accuracy', ascending=False)
```

Out[144]:

	Models	Accuracy	AUC
5	Gradient Boosting	91.751670	0.745367
2	Random Forest Classifier	91.554341	0.729820
1	Decision Tree Classifier	91.314511	0.927086
0	Logistic Model	91.071645	0.688099
3	K-Near Neighbors	90.497875	0.634491
4	Gaussian NB	83.913175	0.736643


```


In [148]: #fig, (ax1, ax2, ax3, ax4, ax5) = plt.subplots(nrows = 2, ncols = 3, figsize =
(15, 4))
fig, ax_arr = plt.subplots( nrows = 1, ncols = 1,figsize = (5,5))

fpr1, tpr1, thresholds = roc_curve(y_test, gbk_pred_proba)
ax_arr.plot(fpr1, tpr1, 'b', label = 'GradientBoosting', color='black')
fpr2, tpr2, thresholds = roc_curve(y_test, gauss_pred_proba)
#fpr1, tpr1, thresholds = roc_curve(y_test, gauss_pred_proba)
#plt.plot(fpr, tpr, Label='GaussianNB (AUC = %0.03f)' % roc_auc)
ax_arr.plot(fpr2, tpr2, 'b', label = 'GaussianNB', color='red')
fpr3, tpr3, thresholds = roc_curve(y_test, knn_pred_proba)
#plt.plot(fpr, tpr, Label='KNN (AUC = %0.03f)' % roc_auc)
ax_arr.plot(fpr3, tpr3, 'b', label = 'Decision Tree', color='blue')
fpr4, tpr4, thresholds = roc_curve(y_test, rd_pred_proba)
#plt.plot(fpr, tpr, Label='RandomForestClassifier (AUC = %0.03f)' % roc_auc)
ax_arr.plot(fpr4, tpr4, 'b', label = 'Knn', color='brown')
fpr5, tpr5, thresholds = roc_curve(y_test, dec_pred_proba)
#plt.plot(fpr, tpr, Label='DecisionTreeClassifier (AUC = %0.03f)' % roc_auc)
ax_arr.plot(fpr5, tpr5, 'b', label = 'Random Forest', color='green')
fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba)
#plt.plot(fpr, tpr, Label='LogisticRegression (AUC = %0.03f)' % roc_auc)
ax_arr.plot(fpr, tpr, 'b', label = 'Logistic', color='grey')
ax_arr.set_title('Receiver Operating Comparison ',fontsize=16)
ax_arr.set_ylabel('True Positive Rate',fontsize=15)
ax_arr.set_xlabel('False Positive Rate',fontsize=15)
ax_arr.legend(loc = 'lower right', prop={'size': 12})

#plt.subplots_adjust(wspace=0.2)
#plt.tight_layout()

```

Out[148]: <matplotlib.legend.Legend at 0x2afe4abb518>




```
In [146]: #df = pd.read_csv(r"C:\Users\bhara\Desktop\PR\FinalData\full_le.csv")
#X = df.drop('y', axis=1).values
#y = df['y'].values
#pp=df.drop('y', axis=1)
#x_train, x_test, y_train, y_test = train_test_split(X, y, test_size = 0.3, random_state=42)
rfc = RandomForestClassifier(n_estimators=100)
rfc.fit(x_train, y_train)
feature_importances = pd.DataFrame(rfc.feature_importances_, index = X.columns,
columns=['importance']).sort_values('importance', ascending=False)
feature_importances
```

Out[146]:

	importance
duration	0.330269
euribor3m	0.114898
age	0.099301
nr.employed	0.051417
campaign	0.044899
day_of_week	0.043404
job	0.039969
education	0.039507
pdays	0.034222
poutcome	0.031343
cons.conf.idx	0.029107
cons.price.idx	0.022261
emp.var.rate	0.022013
housing	0.021167
month	0.019161
marital	0.018103
loan	0.015799
previous	0.012909
contact	0.010250

```
In [147]: ax = feature_importances.plot.barh(rot=0, figsize = (8,8))  
plt.title("Determining Feature importances \n with DecisionTreeClassifier", fo  
ntsize=12)
```

```
Out[147]: Text(0.5, 1.0, 'Determining Feature importances \n with DecisionTreeClassifie  
r')
```



```
In [ ]:
```

```
In [ ]:
```